

Machine Learning

Week 2: Pattern Classification

Maheesan Niranjana

School of Electronics and Computer Science
University of Southampton

Autumn Semester 2016/17

Overview (Week 2)

- Review of what we learnt in Lab One
 - Multivariate Gaussian
 - Drawing samples from $\mathcal{N}(\mathbf{m}, \mathbf{C})$
 - Principal directions
- Introduction to Bayesian Decision Theory
- Bayes' Classifier for Simple Gaussian Distributions
- Simple Classifiers
 - Distance to mean classifier
 - Nearest Neighbour classifier
 - Linear classifier (more on this later)
 - Perceptron (formal setting later)
- What will we learn in Lab Two?

Bayesian Decision Theory

- Classes: $\omega_i, i = 1, \dots, K$
- Prior Probabilities: $P[\omega_1], \dots, P[\omega_K]$;
 $P[\omega_i] \geq 0, \sum_{i=1}^K P[\omega_i] = 1$
- Likelihoods (class conditional probabilities): $p(\mathbf{x}|\omega_i), i = 1, \dots, K$
- Posterior Probability: $P[\omega_j | \mathbf{x}]$

$$P[\omega_j | \mathbf{x}] = \frac{p(\mathbf{x} | \omega_j) P[\omega_j]}{\sum_{i=1}^K p(\mathbf{x} | \omega_i) P[\omega_i]}$$

- From prior knowledge: $P[\omega_i]$; From training data: $p(\mathbf{x}|\omega_i)$
- Decision rule: Assign \mathbf{x} to the class that maximizes posterior probability.
- The denominator is a constant; i.e. does not depend on ω_j
- Hence the decision rule becomes:

$$\mathbf{x} \in \max_j p(\mathbf{x} | \omega_j) P[\omega_j]$$

Bayes' Classifier for Gaussian Densities

Make assumptions, cancel common terms when making comparisons...

- Decision rule from: $p(\mathbf{x} | \omega_j) P[\omega_j]$
- Assume the two classes are Gaussian distributed with distinct means and identical covariance matrices
 $p(\mathbf{x} | \omega_j) = \mathcal{N}(\mathbf{m}_j, \mathbf{C})$
- Substitute into Bayes' classifier decision rule

$$\begin{aligned} P[\omega_1 | \mathbf{x}] &\leq P[\omega_2 | \mathbf{x}] \\ p(\mathbf{x} | \omega_1) P[\omega_1] &\leq p(\mathbf{x} | \omega_2) P[\omega_2] \end{aligned}$$

$$\begin{aligned} \frac{1}{(2\pi)^{p/2}(\det(\mathbf{C}))^{1/2}} \exp \left\{ -\frac{1}{2}(\mathbf{x} - \mathbf{m}_1)^t \mathbf{C}^{-1}(\mathbf{x} - \mathbf{m}_1) \right\} P[\omega_1] &\leq \\ \frac{1}{(2\pi)^{p/2}(\det(\mathbf{C}))^{1/2}} \exp \left\{ -\frac{1}{2}(\mathbf{x} - \mathbf{m}_2)^t \mathbf{C}^{-1}(\mathbf{x} - \mathbf{m}_2) \right\} P[\omega_2] \end{aligned}$$

Bayes' classifier for simple densities (cont'd)

Distinct Means; Equal, isotropic covariance matrix

- Suppose the densities are isotropic and priors are equal
i.e. $\mathbf{C} = \sigma^2 \mathbf{I}$ and $P[\omega_1] = P[\omega_2]$
- The comparison simplifies to (see algebra on board):

$$\begin{aligned}(\mathbf{x} - \mathbf{m}_1)^t (\mathbf{x} - \mathbf{m}_1) &\leq (\mathbf{x} - \mathbf{m}_2)^t (\mathbf{x} - \mathbf{m}_2) \\ |\mathbf{x} - \mathbf{m}_1| &\leq |\mathbf{x} - \mathbf{m}_2|\end{aligned}$$

- The above is a simple *distance to mean* classifier
- Under the above simplistic assumptions, we only need to store one template per class (the means)!

Bayes' classifier for simple densities (cont'd)

Distinct Means; Common covariance matrix (but not isotropic)

- Cancel common terms and take log

$$(\mathbf{x} - \mathbf{m}_1)^t \mathbf{C}^{-1} (\mathbf{x} - \mathbf{m}_1) \leq (\mathbf{x} - \mathbf{m}_2)^t \mathbf{C}^{-1} (\mathbf{x} - \mathbf{m}_2) - \log \left\{ \frac{P[\omega_1]}{P[\omega_2]} \right\}$$

- Also simplifies to a linear classifier!

$$\mathbf{w}^t \mathbf{x} + b \leq 0$$

$$\mathbf{w} = 2\mathbf{C}^{-1} (\mathbf{m}_2 - \mathbf{m}_1)$$

$$b = \left(\mathbf{m}_1^t \mathbf{C}^{-1} \mathbf{m}_1 - \mathbf{m}_2^t \mathbf{C}^{-1} \mathbf{m}_2 \right) - \log \left\{ \frac{P[\omega_1]}{P[\omega_2]} \right\}$$

- Also a distance to template classifier, where the distance is

$$(\mathbf{x} - \mathbf{m}_1)^t \mathbf{C}^{-1} (\mathbf{x} - \mathbf{m}_1)$$

Known as Mahalanobis distance

Implementing a linear classifier: Perceptron

Error correcting learning

Linear classifier

$$\mathbf{w}^t \mathbf{x} + b \leq 0$$

Expand dimensions: $\mathbf{a} = [\mathbf{w}^t \ b]^t$ and $\mathbf{y} = [\mathbf{x}^t \ 1]^t$

$$\mathbf{a}^t \mathbf{y} \leq 0$$

random guess of the weights

repeat

 select data at random

 if not correctly classified

 update weights

until (all data correctly classified)

Update:

$$\mathbf{a}^{(k+1)} = \mathbf{a}^{(k)} + \eta \mathbf{y}^{(k)}$$

Lab 2

- 1 Some plotting
- 2 Bayes' optimal class boundary
- 3 Implement your own perceptron algorithm