Machine Learning

Week 4: Linear Regression & Perceptron

Mahesan Niranjan

School of Electronics and Computer Science
University of Southampton

Autumn Semester 2016/17

Linear Regression & Perceptron

- Data: $\{\boldsymbol{x}_n, f_n\}_{n=1}^N$
  Input: $\boldsymbol{x}_n \in \mathcal{R}^p$; target / output f_n real valued
- Model: $f = \boldsymbol{w}^t \boldsymbol{x} + w_0$
  Output linear function of input (including a constant w_0)
- Work in $(p+1)$ dimensional space to avoid treating w_0 separately

$$\boldsymbol{y} = \begin{pmatrix} \boldsymbol{x} \\ 1 \end{pmatrix} \quad \boldsymbol{a} = \begin{pmatrix} \boldsymbol{w} \\ w_0 \end{pmatrix}$$

- Data: $\{\boldsymbol{y}_n, f_n\}_{n=1}^N$
- Model: $f = \boldsymbol{y}^t \boldsymbol{a}$
- $p+1$ unknowns held in vector $\boldsymbol{a}$

Error and Minimization

- $E = \sum_{n=1}^{N} \left\{ \boldsymbol{y}_n^t \boldsymbol{a} - f_n \right\}^2$

- $E = \sum_{n=1}^{N} \left\{ \left(\sum_{j=1}^{(p+1)} a_j \, y_{nj} \right) - f_n \right\}^2$

- To find the best $\boldsymbol{a}$ we minimize E – differentiate with respect to each of the unknowns in $\boldsymbol{a}$ and set to zero.

-
$$\frac{\partial E}{\partial a_i} = 2 \sum_{n=1}^{N} \left\{ \left(\sum_{j=1}^{(p+1)} a_j \, y_{nj} \right) - f_n \right\} (y_{ni})$$

- There are $(p+1)$ derivatives (with respect to each a_i)
- Equating them to zero gives $(p+1)$ equations in $(p+1)$ unknowns

Solution to Regression

- $(p+1)$ simultaneous equations to solve:
  i^{th} row, j^{th} column shown

$$
\begin{pmatrix}
& \cdots & \\
\cdots & \cdots & \cdots \\
\vdots & \sum_{n=1}^{N} y_{ni} y_{nj} & \cdots \\
\vdots & \cdots & \vdots \\
\cdots & \cdots & \cdots
\end{pmatrix}
\begin{pmatrix}
a_1 \\
a_2 \\
\vdots \\
a_{(p+1)}
\end{pmatrix}
=
\begin{pmatrix}
\vdots \\
\sum_{n=1}^{N} f_n y_{ni} \\
\vdots
\end{pmatrix}
$$

Derivation in vector/matrix form

- $\boldsymbol{Y}$: $N \times (p+1)$ matrix n^{th} row is $\boldsymbol{y}_n^t$
- $\boldsymbol{f}$: $N \times 1$ vector of outputs
- Error $E = ||\boldsymbol{Y}\boldsymbol{a} - \boldsymbol{f}||^2$
- $\boxed{\text{Homework:}}$ Verify the error written like this is the same as the one we wrote out in lengthy algebra.
- Gradient

$$\nabla_{\boldsymbol{a}} E = 2\boldsymbol{Y}^t (\boldsymbol{Y}\boldsymbol{a} - \boldsymbol{f})$$

- Equating the gradient to zero gives

$$\begin{aligned} \boldsymbol{Y}^t\boldsymbol{Y}\boldsymbol{a} &= \boldsymbol{Y}^t\boldsymbol{f} \\ \boldsymbol{a} &= \left(\boldsymbol{Y}^t\boldsymbol{Y}\right)^{-1}\boldsymbol{Y}^t\boldsymbol{f} \end{aligned}$$

- $\boxed{\text{Homework:}}$ With three data points in one dimensional input space (x_1, f_1), (x_2, f_2) and (x_3, f_3) and two unknowns, slope (m) and intercept (c) of fitting a straight line, write out all the expressions seen so far.

Solution by Gradient Descent

- Gradient vector: $\nabla_{\boldsymbol{a}} E = 2\boldsymbol{Y}^t (\boldsymbol{Y}\boldsymbol{a} - \boldsymbol{f})$
- Steepest descent algorithm:

```
Initialize a at random
Update a^(k+1) = a^(k) − η ∇_a E
Until Convergence
```

- Second order (Newton's) method

```
Initialize a at random
Update a^(k+1) = a^(k) − η H^−1 ∇_a E
Until Convergence
```

- Rapid convergence with second order method, but cost of computing and inverting $\boldsymbol{H}$ can be high (more on this under Neural Networks)

Gradient and Stochastic Gradient Descent

- Error $E = \sum_{n=1}^{N} e_n^2$
- True gradient:

$$\nabla_{\boldsymbol{a}} E = 2 \sum_{n=1}^{N} \left\{ \mathbf{y_n^t a} - \mathbf{f_n} \right\} (\mathbf{y_n})$$

- Gradient computed on n^{th} data:

$$\nabla_{\boldsymbol{a}} e_n = 2 \left\{ \mathbf{y_n^t a} - \mathbf{f_n} \right\} (\mathbf{y_n})$$

Regularization

- Pseudo inverse solution: $\boldsymbol{a} = \left(\boldsymbol{Y}^t \boldsymbol{Y} \right)^{-1} \boldsymbol{Y}^t \boldsymbol{f}$
- This can be ill conditioned, so we could *regularize* by

$$\boldsymbol{a} = \left(\boldsymbol{Y}^t \boldsymbol{Y} + \gamma \boldsymbol{I} \right)^{-1} \boldsymbol{Y}^t \boldsymbol{f}$$

  where γ is a small constant.
- We achieve precisely this by minimizing an error of the form

$$|| \boldsymbol{Ya} - \boldsymbol{f} ||^2 + \gamma || \boldsymbol{a} ||^2$$

  Here a quadratic penalty term has been included
- Homework: Differentiate this error and derive the regularized solution
- Sparse solutions are obtained by regularizing with an l_1 norm (sum of absolute values of $\boldsymbol{a}$, i.e. $\sum_{j=1}^{p} |a_j|$); See **Lab 4**.

Perceptron
A suitable performance measure

- Number of misclassified examples as measure of error
  Piecewise constant (cannot differentiate)
- Suitable error measure:

$$E_P = - \sum y_n^t a$$

- Summation taken over misclassified examples
- We started with $y_n^t a > 0$ for positive class and $y_n^t a < 0$ for the negative class; we then switch the signs of negative class examples and required $y_n^t a > 0$ for all the training data; so for the misclassified examples $- \sum y_n^t a$ should be as small as possible.

Perceptron
Learning rule

- Gradient:

$$\frac{\partial E}{\partial a} = - \sum y_n$$

- Gradient algorithm: $a^{(k+1)} = a^{(k)} + \sum y_n$
- Stochastic gradient algorithm:

$$a^{(k+1)} = a^{(k)} + y_n$$

- Note what y_n is. It is an item of data that is taken at random and happens to be misclassified by the current value of a at iteration k.

Perceptron
Convergence of the learning rule

- Learning Rule: $\boldsymbol{a}^{(k+1)} = \boldsymbol{a}^{(k)} + \boldsymbol{y}(k)$
  where $\boldsymbol{y}(k)$ is a misclassified input.
- Training criterion
  - We start with requiring $\boldsymbol{a}^t \boldsymbol{y}(k) \lessgtr 0$, depending on the example belonging to class 1 or class 2.
  - If we switch the signs of examples of class 2, we require $\boldsymbol{a}^t \boldsymbol{y}(k) > 0$ for all k.
- On misclassified data $\boldsymbol{a}^t \boldsymbol{y}(k) < 0$
- If $\widehat{\boldsymbol{a}}$ is a solution (separable data), for all k, $\widehat{\boldsymbol{a}} \boldsymbol{y}(k) > 0$
- We prove convergence by showing:
  $||\boldsymbol{a}^{(k+1)} - \widehat{\boldsymbol{a}}||^2 < ||\boldsymbol{a}^{(k)} - \widehat{\boldsymbol{a}}||^2$ for this update rule. *i.e.* the learning rule brings the guess closer to a valid solution.

Perceptron
Convergence of the learning rule (cont'd)

- For perceptron criterion, the magnitude of $\boldsymbol{a}$ is not relevant (only the direction is). Hence for some scalar α, we wish to show

$$||\boldsymbol{a}^{(k+1)} - \alpha\,\widehat{\boldsymbol{a}}||^2 < ||\boldsymbol{a}^{(k)} - \alpha\,\widehat{\boldsymbol{a}}||^2$$

- From the update formula
$$\boldsymbol{a}^{(k+1)} - \alpha\,\widehat{\boldsymbol{a}} = \boldsymbol{a}^{(k)} - \alpha\,\widehat{\boldsymbol{a}} + \boldsymbol{y}(k)$$

- Taking magnitudes
$$||\boldsymbol{a}^{(k+1)} - \alpha\widehat{\boldsymbol{a}}||^2 = ||\boldsymbol{a}^{(k)} - \alpha\widehat{\boldsymbol{a}}||^2 + 2(\boldsymbol{a}^{(k)} - \alpha\widehat{\boldsymbol{a}})^t \boldsymbol{y}(k) + ||\boldsymbol{y}(k)||^2$$

- If we drop the negative term $\boldsymbol{a}^{(k)^t} \boldsymbol{y}(k)$ from RHS, the equality becomes an inequality
$$||\boldsymbol{a}^{(k+1)} - \alpha\widehat{\boldsymbol{a}}||^2 < ||\boldsymbol{a}^{(k)} - \alpha\widehat{\boldsymbol{a}}||^2 - 2\alpha\widehat{\boldsymbol{a}}^t \boldsymbol{y}(k) + ||\boldsymbol{y}(k)||^2$$

Perceptron
Convergence of the learning rule (cont'd)

- Of the three terms on the right hand side, we know $\widehat{\boldsymbol{a}}^t \boldsymbol{y}(k) > 0$, because $\widehat{\boldsymbol{a}}$ is assumed to be a solution.
- If we select

$$
\begin{aligned}
\beta^2 &= \max_i \|\boldsymbol{y}_i\|^2 \\
\gamma &= \min_i \widehat{\boldsymbol{a}}^t \boldsymbol{y}_i
\end{aligned}
$$

i.e. largest of the positive term and smallest of the negative term, then for $\alpha = \beta^2/\gamma$,

$$
\|\boldsymbol{a}^{(k+1)} - \alpha\widehat{\boldsymbol{a}}\|^2 < \|\boldsymbol{a}^{(k)} - \alpha\widehat{\boldsymbol{a}}\|^2 - \beta^2
$$

- (Note the inequality remains true when the right hand side is replaced by a quantity larger than what it previously was.)
- Every correction takes the guess closer to a true solution.
- From an initialization $\boldsymbol{a}^{(1)}$, we will find a solution in *at most* $k_0 \frac{\|\boldsymbol{a}(1) - \alpha\widehat{\boldsymbol{a}}\|^2}{\beta^2}$ updates.

Summary

- Linear regression
  - Solution as pseudo inverse
  - Solution by gradient descent
  - Regularization
- Perceptron
  - Setting up a suitable error function
  - Convergence of the algorithm